

MVA 2008 - 2009

Object recognition and computer vision

Assignment 2: Stitching photo mosaics

Jean Ponce, Cordelia Schmid and Josef Sivic

Due date: November 18th 2008

(adapted from Alyosha Efros, CMU, Svetlana Lazebnik, UNC, and Andrew Zisserman, Oxford)

The goal of the assignment is to automatically stitch images acquired by a panning camera into a mosaic as illustrated in figures 1 and 2 below.



Fig.1: Three images acquired by a panning camera.



Fig.2: Images stitched to a mosaic.

Algorithm outline:

1. Choose one image as the reference frame.
2. Estimate homography between each of the remaining images and the reference image. To estimate homography between two images use the following procedure:
 - a. Detect local features in each image.

- b. Extract feature descriptor for each feature point.
 - c. Match feature descriptors between two images.
 - d. Robustly estimate homography using RANSAC.
3. Warp each image into the reference frame and composite warped images into a single mosaic.

Tips and detailed description of the algorithm:

The algorithm will be described on the example of stitching images shown in figure 1.

1. **Choosing the reference image.** Choosing the middle image of the sequence as the reference frame is the preferred option, as it will result in mosaics with less distortion. However, you can choose any image of the sequence as the reference frame. If you choose image 1 or 3 as the reference frame, directly estimating homography between images 1 and 3, will be difficult as they have very small overlap. To deal with this issue, you might have to “chain” homographies, e.g. $H_{13} = H_{12} * H_{23}$.
2. **Estimating homography.**
 - a. **Detect local features in each image.** Use your feature detector from assignment 1. You can also use function [harris.m](#) by A. Efros implementing a simple single scale Harris corner detector.
 - b. **Extract feature descriptor for each feature point.** Implement feature descriptor extraction outlined in Section 4 of the paper [“Multi-Image Matching using Multi-Scale Oriented Patches” by Brown et al.](#) Don’t worry about rotation-invariance – just extract axis-aligned 8x8 patches. Note that it’s important to sample these patches from the larger 40x40 window to have a nice large blurred descriptor. Don’t forget to bias/gain-normalize the descriptors. Ignore the Haar wavelet transform section.
 - c. **Match feature descriptors between two images.** Implement Feature Matching (Section 5 in [“Multi-Image Matching using Multi-Scale Oriented Patches” by Brown et al.](#)). That is, you will need to find pairs of features that look similar and are thus likely to be in correspondence. We will call this set of matched features “tentative” correspondences. You may find function [dist2.m](#) useful for distance computations. For thresholding, use the simpler approach due to Lowe of thresholding on the ratio between the first and the second nearest neighbors. Consult Figure 6b in the paper for picking the threshold. Ignore Section 6 on fast Wavelet indexing. You can visualize the tentative correspondences between two images by displaying the feature displacements. For example, to visualize tentative correspondences between image 1 and 2: (i) show image 1, (ii) show detected features in image 1 (display only region centers as points, do not worry about the regions’ scale), (iii) show displacements between detected features in image 1 and matched features in image 2 by line segments. This is illustrated in figure 3 and can be achieved using the following matlab code, where `im1_pts` and `im2_pts` are 2-by-n matrices holdin (x,y) image locations of tentatively corresponding features in image 1 and image 2, respectively:

```
figure; clf; imagesc(imlrgb); hold on;

% show features detected in image 1
plot(im1_pts(1,:),im1_pts(2,:),'+g');

% show displacements
line([im1_pts(1,:); im2_pts(1,:)],[im1_pts(2,:); im2_pts(2,:)],'color','y')
```
 - d. **Robustly estimate homography using RANSAC.** Use a sample of 4-points to compute each homography hypothesis. You will need to write a function of the form:

$$H = \text{computeH}(\text{im1_pts}, \text{im2_pts})$$

where, again, `im1_pts` and `im2_pts` are 2-by-n matrices holding the (x,y) locations of $n(=4)$ point correspondences from the two images and H is the recovered 3x3 homography matrix. In order to compute the entries in the matrix H , you will need to set up a linear system of n equations (i.e. a matrix equation of

the form $Ah=0$ where h is a vector holding the 8 unknown entries of H). The solution to the homogeneous least squares system $Ax=0$ is obtained from the SVD of A by the singular vector corresponding to the smallest singular value. In Matlab: `[U,S,V]=svd(A); x = V(:,end);` For more details on homography estimation from point correspondences see a [note written by David Kriegman](#). Don't forget to re-estimate the homography from all inlying points.

For RANSAC, a very simple implementation performing a fixed number of sampling iterations is sufficient. You should output a single transformation that gets the most inliers in the course of all the iterations. For the various RANSAC parameters (number of iterations, inlier threshold), play around with a few "reasonable" values and pick the ones that work best. For randomly sampling matches, you can use the Matlab `randperm` function. You should display the set of inliers as illustrated in figure 3.

3. **Warping and compositing.** Warp each image into the reference frame using the estimated homography and composite warped images into a single mosaic. You can use the [vgg_warp_H](#) function for warping. Here is an example code to warp and composite images 1 and 2 :

First, define a mosaic image to warp all the images onto. Here we assume image 2 as the reference image, and map this image to the origin of the mosaic image using the identity homography (`eye(3)` in Matlab).

```
bbox = [-400 1200 -200 700] % image space for mosaic
Im2w = vgg_warp_H(Im2, eye(3), 'linear', bbox); % warp image 1 to mosaic image
```

Now warp image 1 to a separate mosaic image using estimated homography H_{12} between image 1 and image 2

```
Im1w = vgg_warp_H(Im1, H12, 'linear', bbox);
```

and finally combine the mosaic images by taking the pixel with maximum value from each image. This tends to produce less artefacts than taking the average of warped images.

```
imagesc(double(max(Im1w, Im2w))/255);
```




Fig. 3: Visualization of tentative correspondences (left) and homography inliers (right). Detected points are shown in green. Feature displacements are shown by yellow lines.

Test images

Here are the [three images](#) from figure 1. Also run your code on a new set of at least two homography related images, which you capture yourself using a digital camera or download from the Internet. To capture the images, shoot from the same point of view but with different view directions, and with overlapping fields of view. Make sure the images have enough overlap (e.g. at least 40%). Another way to do this is to shoot pictures of a planar surface (e.g. a wall) from different points of view.

What to hand in

You should prepare a (very brief) report including the following:

- The final mosaic of the provided test images (similar to figure 2).
- Visualization of the set of tentative matches and the set of inliers for one image pair (similar to figure 3).
- Give the parameter values you have used (e.g. the threshold for descriptor matching, number of iterations and the inlier threshold for RANSAC).
- Show the input images and the final mosaic for your new image set (see the Test images section above).

Send your report, code, and your additional test images in a single zip file to **Josef Sivic** <Josef.Sivic@ens.fr>.

Helpful resources

- Nice [Slides](#) by Alyosha Efros on mosaicing.
- A note on [computing homography](#) by David Kriegman.
- Another note on [computing Homography](#).
- [Tutorial on image alignment and stitching \(draft\)](#) by Richard Szeliski.
- M. Brown, R. Szeliski and S. Winder. [Multi-Image Matching using Multi-Scale Oriented Patches](#). *International Conference on Computer Vision and Pattern Recognition (CVPR2005)*, pages 510-517
- A comprehensive treatment of homography estimation can be found in chapter 4 of Multiple View Geometry in Computer Vision by R. Hartley and A. Zisserman, Cambridge University Press, 2004.